

Feature Request 1

#597 [Performance] Fibonacci Heaps for Ball Tree

Link: <https://github.com/scikit-learn/scikit-learn/issues/597>

Description:

The Ball Tree algorithm is a method of performing the Nearest Neighbor search. Current implementation relies on Dijkstra's algorithm which achieves its peak theoretical speed with a Fibonacci heap. So Ball Tree Nearest Neighbor search could potentially be sped up using Fibonacci heaps instead of the current implementation.

- an implementation of Fibonacci heap already exists in the `sklearn/utils/graph_nearest_neighbor.pyx` cython file.
- It would be refactored to fit requirements, and be called whenever Dijkstra's algorithm is used. This solution would "catch" all instances of different classification algorithms involving BallTree.
- Dijkstra's algorithm is only used in the `graph_shortest_path.pyx` file.

Changes:

BallTree itself requires no change, it is only when performing the Nearest Neighbor search on a constructed BallTree instance, that we want to change the data structures used.

In the `graph.py` utility file, `graph_shortest_path` is imported:

```
15 from .graph_shortest_path import graph_shortest_path # noqa
```

The implementation of Dijkstra's algorithm regarding `graph_shortest_path` is found in the `graph_shortest_path.pyx`

File: https://github.com/scikit-learn/scikit-learn/blob/main/sklearn/utils/graph_shortest_path.pyx

In this file we want to change the current implementation of `min_heap` into `fib_heap`.

```
30 def graph_shortest_path(dist_matrix, directed=True, method='auto'):  
31     """  
32     Perform a shortest-path graph search on a positive directed or  
33     undirected graph.  
34
```

When successfully implemented, the amortized runtime of Nearest Neighbor should be lowered.

```
>>> nbrs = NearestNeighbors(n_neighbors=2, algorithm='ball_tree').fit(X)  
>>> distances, indices = nbrs.kneighbors(X)
```

Diagram:

All implementation methods are within the class of graph_shortest_path

